

Dr. Sarah Chen

Email: sarah.chen@aimail.com

Phone: (555) 123-4567 | LinkedIn: linkedin.com/in/sarahchen

GitHub: github.com/sarahchen-ml | Location: San Francisco, CA

Professional Summary

Senior Machine Learning Engineer with 8+ years of experience in deep learning, natural language processing, and computer vision. Led teams of 5-10 engineers. Published 12 papers in top-tier conferences (NeurIPS, ICML, ACL).

Passionate about building production ML systems that scale.

Experience

Senior ML Engineer | DataCorp Inc. | Jan 2022 - Present

Leading the recommendation system team, serving 10M+ daily active users.

- Designed and deployed real-time recommendation engine using Python, TensorFlow, and Apache Kafka, improving click-through rate by 23%
- Built A/B testing framework that reduced experiment cycle time from 2 weeks to 3 days
- Mentored 4 junior engineers and conducted 50+ technical interviews
- Implemented automated model monitoring with Prometheus and Grafana
- Technologies: Python, TensorFlow, PyTorch, Spark, AWS, Docker, Kubernetes

ML Engineer | AI Labs Research | Jun 2019 - Dec 2021

Developed NLP models for document understanding and information extraction.

- Created BERT-based model for legal document classification (95% accuracy)
- Built named entity recognition system for medical records using spaCy
- Optimized model inference latency by 40% using TensorRT optimization
- Published 3 papers at ACL and EMNLP conferences
- Technologies: Python, PyTorch, Hugging Face, spaCy, Java, GCP

Data Scientist | TechStart AI | Jul 2017 - May 2019

Built predictive models for customer churn and lifetime value estimation.

- Developed XGBoost model predicting customer churn with 89% AUC
- Created dashboard for real-time business metrics using React and D3.js
- Automated data pipeline reducing manual reporting by 70%
- Technologies: Python, R, SQL, scikit-learn, Tableau, AWS Redshift

Education

Ph.D. in Computer Science (Machine Learning)

Stanford University | 2017

Dissertation: Scalable Methods for Distributed Deep Learning

Advisor: Prof. Andrew Ng

B.S. in Computer Science & Mathematics

MIT | 2013

Summa Cum Laude, GPA: 3.95/4.0

Technical Skills

Programming Languages: Python, Java, C++, R, SQL, JavaScript, Go

ML Frameworks: TensorFlow, PyTorch, Keras, scikit-learn, XGBoost, LightGBM

NLP Tools: Hugging Face Transformers, spaCy, NLTK, Gensim

Cloud Platforms: AWS (SageMaker, EC2, S3), GCP (Vertex AI, BigQuery), Azure ML

DevOps: Docker, Kubernetes, CI/CD, Git, Terraform, Prometheus, Grafana

Databases: PostgreSQL, MongoDB, Redis, Elasticsearch, Neo4j

Selected Publications

1. Chen et al. (2024) 'Scalable Distributed Training of Large Language Models'
NeurIPS 2024
2. Chen, S. et al. (2023) 'Efficient Fine-tuning for Domain-Specific NLP'
ICML 2023
3. Chen et al. (2022) 'Real-time Recommendation Systems at Scale'
RecSys 2022
4. Chen, S. (2021) 'Legal Document Classification with BERT'
ACL 2021
5. Chen et al. (2020) 'Medical NER with Transfer Learning'
EMNLP 2020

Certifications

- AWS Machine Learning Specialty (2023)
- Google Cloud Professional ML Engineer (2022)
- TensorFlow Developer Certificate (2021)

Languages

English (native), Mandarin (fluent), Spanish (conversational)